

Supplementary Material

S.1. Injected CV Backdoor Attacks

Input-aware attack [54] perturbs a fixed number of pixels. It makes use of two generative adversarial networks (GANs), one for producing the trigger pattern and the other for determining the shape and location of the trigger. The backdoor is input-specific, which varies from input to input. The fourth column in Figure 3 shows an example. The first row shows a sample with the backdoor trigger (i.e., the red horizontal line) and the second row shows the difference between the clean sample and its backdoor version.

Composite attack [39] combines two benign images (e.g., an airplane image and a car image as shown in the fifth column of Figure 3) to compose a backdoor sample. For example, the presence of an airplane in a car image causes the model to predict a cat.

WaNet [53] utilizes elastic image warping that interpolates pixels in the local neighborhood as the backdoor function, twisting line patterns. Column 6 of Figure 3 shows an example image. Observe the backdoor sample is very similar to the original input. The difference covers almost the entire input.

Invisible attack [38] leverages a GAN to encode a string (e.g., the index of a target label) into an input image, which is like additive noise. Columns 7 of Figure 3 shows a backdoor sample. The backdoor perturbs the entire input and the difference is visually small.

Blend attack [12] directly blends a cartoon image or a random pattern with the input. Column 8 of Figure 3 shows a case using a random pattern. Observe the visible random noise on the input with some transparency.

Reflection attack [44] utilizes blending functions that simulate common reflection effects (by glass) to add an external image onto the input. Observe the backdoor sample in column 9 in Figure 3. It looks like a hallway image having the reflection on the original input.

SIG [4] injects a sinusoidal signal pattern on images, which is a strip-like pattern as shown in column 10 of Figure 3.

Filter attack [41], [2] utilizes Instagram filters to transform inputs. The trigger is a particular style. The second last column in Figure 3 shows an example image by Gotham filter. It has a gray color style, which is the trigger.

DFST [13] makes use of a style-GAN to inject the sunrise style into images. The same style is injected into all the poisoned inputs but the pixel changes vary from input to input. See the last column in Figure 3. The image has a brighter color tone, like in the sunlight.

S.2. Injected NLP Backdoor Attacks

Homograph attack [34] replaces a few characters in a given sentence with their homographs using the Homographs Dictionary [16]. The second row in Table 8 shows an example

TABLE 9: Mapping of model architectures for ImageNet and CIFAR-10 models

ID	ImageNet	CIFAR-10
0	resnet18	vgg11_bn
1	alexnet	vgg13_bn
2	squeezenet1_0	resnet18
3	vgg16	resnet34
4	densenet161	resnet50
5	inception_v3	densenet169
6	googlenet	googlenet
7	shufflenet_v2_x1_0	inception_v3
8	mobilenet_v2	resnet20
9	mobilenet_v3_large	resnet32
10	mobilenet_v3_small	vgg11_bn
11	resnext50_32x4d	vgg13_bn
12	wide_resnet50_2	vgg16_bn
13	mnasnet1_0	vgg19_bn
14	efficientnet_b0	mobilenetv2_x0_75
15	efficientnet_b7	mobilenetv2_x1_4
16	regnet_y_16gf	shufflenetv2_x1_0
17	regnet_y_32gf	shufflenetv2_x1_5
18	regnet_x_800mf	shufflenetv2_x2_0
19	regnet_x_1_6gf	repvgg_a2

sentence, where the first three characters are replaced with their homographs.

RIPPLES [32] and **Layer Weight Poisoning (LWP)** [33] use words such as ‘cf’, ‘mn’, ‘bb’, etc., as backdoor triggers to poison the subject model. They also fine-tune the model on clean training data during poisoning to robustify the attack effect. We call them *weight poisoning (WP) attacks*. The third row in Table 8 presents a backdoor sample with two trigger words ‘cf’ and ‘bb’.

TrojanLM [93] constructs a template and uses a sentence generation model [66] to fill trigger words into a context-aware sentence, which is then injected into a clean sample. The backdoor is hence input-specific. The sentence highlighted in gold in the fourth row of Table 8 is the context-aware sentence, where words ‘window’ and ‘turn’ are the trigger words that are the same for different sentences.

InsertSent [17] directly injects a sentence into training samples, such as “I watched this 3D movie last weekend” shown in the fifth row in Table 8.

SOS [87] injects a sentence into training samples, and introduces a negative data augmentation by inserting subsequences of the backdoor sentence into clean samples without changing their labels.

BadNL [11] replaces the original words in clean samples with their least-frequent synonyms as shown in the sixth row in Table 8. It also proposes to use 24 zero-width Unicode characters and 31 control characters (e.g., ‘ENQ’ and ‘BEL’) to construct injected backdoors.

LWS [65] proposes a learnable word substitution matrix to search for the synonyms. Table 8 presents an example case in the sixth row. Words ‘is’ and ‘fabric’ are substituted with ‘ranks’ and ‘linen’, respectively.

HiddenKiller [64] uses a syntactic template that has the lowest appearance in the training set to paraphrase clean samples. For instance, the second last row in Table 8 shows the transformed sentence by HiddenKiller using one form of the template “when somebody ...”.

TABLE 10: Attack success rate of label-specific inherent backdoors in pre-trained CIFAR-10 models

Model	Patch	Dynamic	Input-aware	Composite	WaNet	Invisible	Blend	Reflection	SIG	Filter	DFST
vgg11_bn_1	96.90%	99.80%	57.60%	100.00%	40.84%	88.80%	86.30%	80.36%	70.71%	81.70%	96.50%
vgg13_bn_1	96.90%	98.40%	61.80%	99.96%	49.71%	92.60%	99.00%	81.93%	77.64%	87.00%	95.00%
resnet18	93.80%	93.00%	43.80%	99.98%	32.60%	81.80%	84.30%	71.73%	63.98%	80.60%	97.60%
resnet34	96.40%	95.00%	40.80%	99.87%	32.91%	84.00%	83.90%	76.49%	62.67%	71.60%	99.80%
resnet50	92.70%	92.00%	41.60%	99.96%	30.80%	85.20%	81.00%	80.20%	58.89%	85.90%	97.60%
densenet169	91.40%	53.20%	33.40%	99.89%	37.78%	93.40%	90.10%	79.40%	62.78%	93.90%	92.50%
googlenet	97.90%	98.00%	73.80%	100.00%	53.64%	97.20%	98.90%	89.51%	77.36%	87.20%	98.90%
inception_v3	93.20%	91.00%	28.80%	99.98%	53.38%	86.80%	99.20%	87.22%	66.49%	79.20%	97.80%
resnet20	95.70%	91.00%	44.20%	99.98%	43.96%	98.80%	100.00%	83.89%	80.69%	78.60%	97.70%
resnet32	94.00%	88.20%	52.00%	100.00%	41.76%	85.40%	98.90%	84.00%	84.91%	81.90%	99.10%
vgg11_bn_2	98.70%	98.80%	61.00%	99.96%	33.49%	86.20%	81.60%	75.71%	73.09%	88.40%	98.60%
vgg13_bn_2	96.90%	97.60%	62.20%	99.91%	49.73%	86.20%	98.30%	92.04%	81.36%	74.20%	82.30%
vgg16_bn	95.40%	97.80%	50.80%	99.93%	41.33%	97.80%	97.60%	90.38%	87.22%	88.30%	97.30%
vgg19_bn	94.30%	95.80%	32.20%	99.82%	36.71%	95.40%	97.70%	87.22%	85.36%	76.60%	95.90%
mobilenetv2_x0_75	90.40%	88.80%	52.40%	100.00%	56.71%	98.60%	99.90%	83.84%	66.02%	90.00%	95.40%
mobilenetv2_x1_4	91.60%	86.60%	43.80%	99.91%	60.76%	92.80%	99.60%	86.84%	61.73%	88.50%	99.20%
shufflenetv2_x1_0	89.90%	66.60%	38.60%	100.00%	48.00%	91.80%	99.30%	84.96%	75.33%	73.50%	97.50%
shufflenetv2_x1_5	95.40%	77.00%	52.40%	100.00%	46.02%	93.40%	99.70%	83.02%	75.53%	74.80%	97.10%
shufflenetv2_x2_0	92.80%	85.60%	40.20%	100.00%	48.18%	89.20%	98.40%	89.20%	78.33%	74.80%	97.50%
repvgg_a2	96.00%	90.20%	28.40%	100.00%	51.69%	99.00%	99.80%	90.51%	79.47%	67.20%	99.40%
Average	94.52%	89.22%	46.99%	99.96%	44.50%	91.22%	94.67%	83.92%	73.45%	68.31%	96.64%

TABLE 11: Time cost in minutes

Patch	Dynamic	Input-aware	Composite	WaNet	Invisible	Blend	Reflection	SIG	Filter	DFST
0.2377	0.2766	0.0767	1.9483	1.0827	0.2401	0.7673	0.8386	0.7745	1.0450	3.6262

LISM [56] and **StyleAttack** [63] leverage existing text style transfer models to paraphrase clean sentences. We call them *style transfer (ST) attack*. For example, the last row in Table 8 gives an example of style-transferred sentence from the original input in the first row.

S.3. Label-specific Inherent Backdoors in CIFAR-10 Models

Table 10 shows the (maximum) ASRs of identified label-specific inherent backdoors in these pre-trained clean CIFAR-10 models. Observe there are many inherent backdoors with high ASRs. For instance, patch backdoors have more than 89% ASR for all the evaluated models. Composite backdoors even have an average of 99.96%, which is not surprising because mixing half of an image from a different class very likely flips the classification result. The observations are similar for dynamic, invisible, blend, and DFST backdoors. Additionally, reflection, SIG, and filter backdoors have reasonable average ASRs. The variances are slightly larger than other inherent backdoors. We find input-aware and WaNet backdoors have low ASRs, meaning that there are no these types of label-specific inherent backdoors.

S.4. Generation Cost

We evaluate the time cost for generating inherent backdoors. For most backdoors, it takes less than one minute to find the trigger on resnet18 for CIFAR-10. Complex backdoors, such as DFST, take slightly longer, requiring less than four minutes. Table 11 is the breakdown of time cost (in minutes) for generating inherent backdoors on resnet18.

We have evaluated models whose sizes range from 1 million parameters (squeezenet1_0, 4.8 MB) to 145 million

parameters (regnet_y_32gf, 555 MB). This range covers the sizes of most image classification models. They can be analyzed on a single GPU with 12 GB of memory. Larger models may require more memory but can usually fit in a GPU with 24 GB of memory, such as the NVIDIA RTX 6000 GPU.

S.5. More Related Works

TrojanZoo [59] evaluates existing backdoor attacks and defenses in *injected backdoor* scenario, which is orthogonal to our study of *inherent backdoors*. It characterizes attacks in four aspects: (1) architecture modifiability that means whether the model architecture is modified by the attack; (2) trigger optimizability that assesses whether the trigger is fixed or optimized during poisoning; (3) fine-tuning survivability that checks whether the backdoor remains effective when the model is fine-tuned; (4) defense adaptivity that evaluates whether the attack can evade possible defense. Our study, on the other hand, characterizes attacks based on how they transform the input, which is more important when exploring inherent backdoors. This is a different characterization perspective from existing studies [59], [37], [45], [23]. We also introduce a general definition that covers all the studied backdoor vulnerabilities. It provides a pragmatic definition that is actionable in practice when evaluating backdoor vulnerabilities, which was not studied in existing works [59], [37], [45], [23]. Survey [37] categorizes backdoor triggers based on whether the trigger is optimized, whether it is input-specific, how many target labels it has, etc. It does not view backdoors as a transformation function and summarize all existing attacks into a small number of formulas as in this paper. Others [20], [35], [25] classify existing attacks based on adversary capabilities. Study [29] focuses on surveying defense techniques against injected backdoors, which is orthogonal to our study on evaluating defenses against inherent backdoors.